


End-to-End vs. Link-by-Link

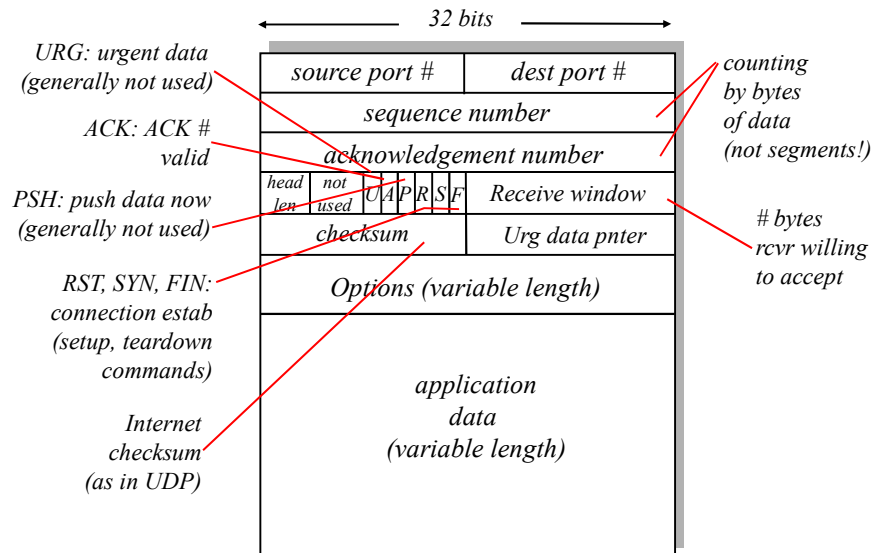
- Can end-to-end reliability be implemented through a link-by-link mechanisms?
- If so, why do we need end-to-end protocols?
- If not, then why do we need link-by-link functionality?



TCP: Overview

- End-to-end issues for providing end-to-end reliable and ordered data transfer
 - Logical connection between two remote hosts
 - RTT is not fixed, even during a connection
 - Packet re-ordering is an issue
 - MSL: Maximum Segment Lifetime
 - Resource discovery
 - Available bandwidth
 - Buffer space
 - Congestion issues in the network

TCP segment format



TCP services/components

- Reliability
 - Sequence numbers and ACKs
 - Time out mechanism
- Flow control
- Connection management
- Congestion control

TCP seq. numbers, ACKs

sequence numbers:

- byte stream “number” of first byte in segment’s data

acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- A: TCP spec doesn’t say, - up to implementor

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



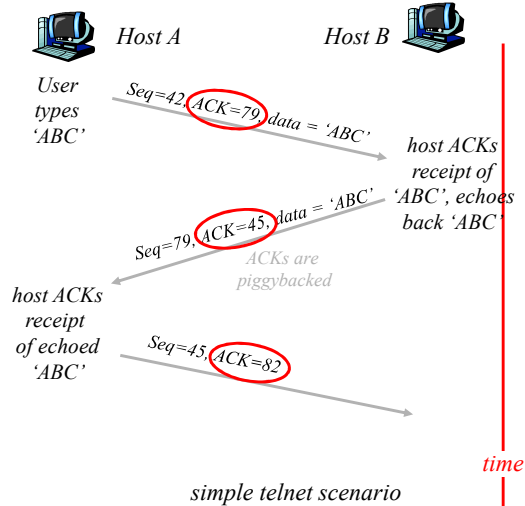
TCP seq. #'s and ACKs

Seq. #'s:

- byte stream “number” of first byte in segment’s data

ACKs:

- seq # of next byte expected from other side
- cumulative ACK



TCP Round Trip Time and Timeout

Q: how to set TCP timeout value?

- longer than RTT
 - but RTT varies
- too short: premature timeout
 - unnecessary retransmissions
- too long: slow reaction to segment loss


Q: how to estimate RTT?

- **SampleRTT**: measured time from segment transmission until ACK receipt
 - ignore retransmissions (Karn/Partridge algorithm)
- **SampleRTT** will vary, want estimated RTT “smoother”
 - average several recent measurements, not just current **SampleRTT**

TCP Round Trip Time, Timeout

$$EstimatedRTT = (1 - \alpha) * EstimatedRTT + \alpha * SampleRTT$$

- ❖ exponential weighted moving average
- ❖ influence of past sample decreases exponentially fast
- ❖ typical value: $\alpha = 0.125$



9


Adaptive Retransmission (The Original Algorithm)

- ❑ Measure **SampleRTT** for each segment/ ACK pair
- ❑ Compute weighted average of RTT
 - $EstRTT = \alpha \times EstRTT + (1 - \alpha) \times SampleRTT$
 - Typically, $0.8 \leq \alpha \leq 0.9$
- ❑ Set timeout based on **EstRTT**
 - $TimeOut = 2 \times EstRTT$

Jacobson/ Karels Algorithm

- Consider variance when setting timeout value
- New Calculations for average RTT
 - $\text{EstRTT} = \delta \times \text{SampleRTT} + (1 - \delta) \times \text{EstRTT}$
 - $\text{Diff} = \text{SampleRTT} - \text{EstRTT}$
 - $\text{Dev} = \delta \times |\text{Diff}| + (1 - \delta) \times \text{Dev}$
 - where δ is a factor between 0 and 1
- $\text{TimeOut} = \mu \times \text{EstRTT} + \phi \times \text{Dev}$
 - where $\mu = 1$ and $\phi = 4$
- Notes
 - accurate timeout mechanism important to congestion control (more about this later)

Karn/Partridge Algorithm



- Do not sample RTT when retransmitting
- Double timeout after each retransmission (exponential Backoff)

TCP reliable data transfer

- TCP provides reliable data transfer service on top of IP's unreliable service
- Cumulative ACKs
- TCP uses single retransmission timer
- Retransmissions are triggered by:
 - timeout events
 - duplicate ACKs
- Initially consider simplified TCP sender:
 - ignore duplicate ACKs
 - ignore flow control, congestion control

TCP sender events

When data is rcvd from application:

- Create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running (think of timer as for oldest unacked segment)
- expiration interval: TimeoutInterval

Upon timeout:

- retransmit segment that caused timeout
- restart timer

When ACK is rcvd:

- If acknowledges previously unACKed segments
 - update what is known to be ACKed
 - start timer if there are still outstanding segments

```

NextSeqNum = InitialSeqNum
SendBase = InitialSeqNum

```

```

loop (forever) {
  switch(event)

```

```

    event: data received from application above
           create TCP segment with sequence number NextSeqNum
           if (timer currently not running)
               start timer
           pass segment to IP
           NextSeqNum = NextSeqNum + length(data)

```

```

    event: timer timeout
           retransmit not-yet-acknowledged segment with
           smallest sequence number
           start timer

```

```

    event: ACK received, with ACK field value of y
           if (y > SendBase) {
               SendBase = y
               if (there are currently not-yet-acknowledged segments)
                   start timer
           }

```

```

} /* end of loop forever */

```

TCP sender (simplified)

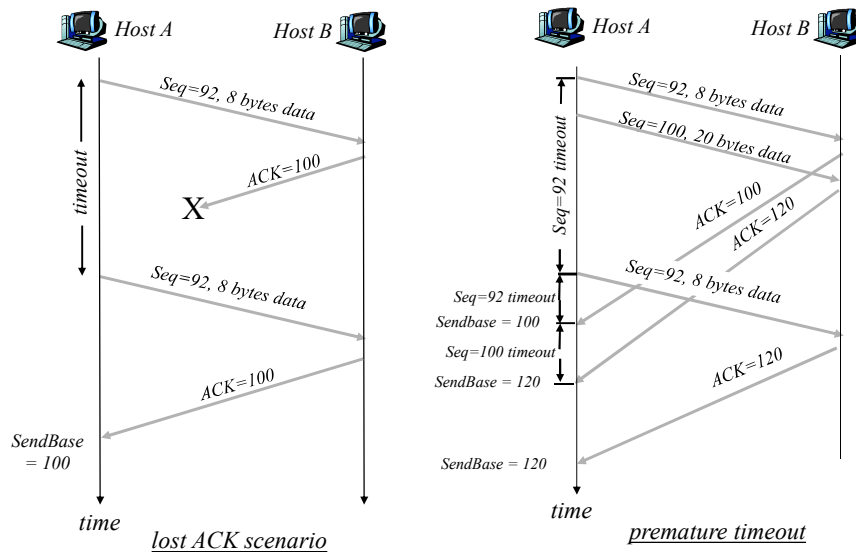
Comment:

- $SendBase-1$: last cumulatively ACKed byte


Example:

- $SendBase-1 = 71$;
 $y = 73$, so the rcvr wants 73+ ;
 $y > SendBase$, so that new data is ACKed

TCP transmission scenarios



TCP transmission scenarios (more)



Wireless Networks Laboratory

Copyright ©by Zygmunt J. Haas, 2017

17

TCP ACK generation [RFC 1122, RFC 2581]

Event at Receiver	TCP Receiver action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK, ACKing in-order segments
Arrival of out-of-order segment higher-than-expected seq. # . Gap detected	Immediately send duplicate ACK, indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap

Wireless Networks Laboratory

Copyright ©by Zygmunt J. Haas, 2017

18

Fast Retransmit

- Time-out period often relatively long:
 - long delay before resending lost packet
- Detect lost segments via duplicate ACKs.
 - Sender often sends many segments back-to-back
 - If segment is lost, there will likely be many duplicate ACKs.
- If sender receives 3 ACKs for the same data, it supposes that segment after ACKed data was lost:
 - **fast retransmit**: resend segment before timer expires

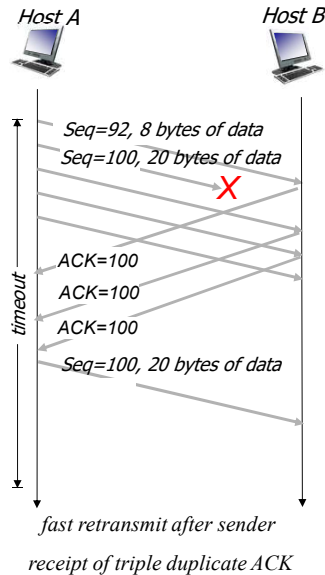
The Fast Retransmit Algorithm:

```
event: ACK received, with ACK field value of y
  if (y > SendBase) {
    SendBase = y
    if (there are currently not-yet-acknowledged segments)
      start timer
  }
  else {
    increment count of dup ACKs received for y
    if (count of dup ACKs received for y = 3) {
      resend segment with sequence number y
    }
  }
```

a duplicate ACK for
already ACKed segment

fast retransmit

TCP Fast Retransmit




21

TCP Flow Control

- receive side of TCP connection has a receive buffer:

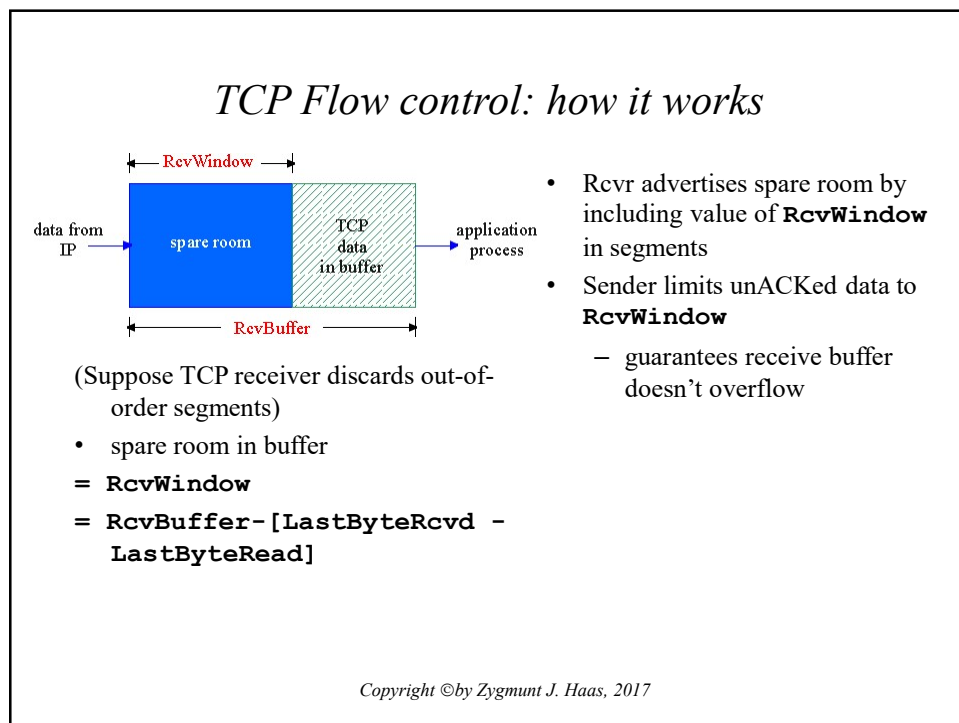
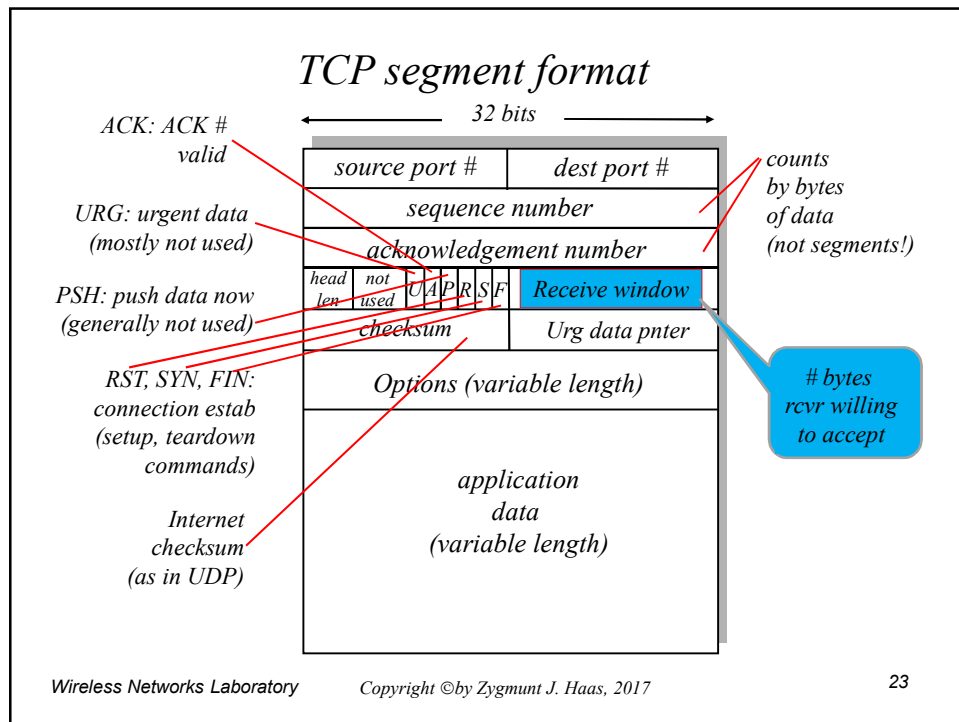
flow control

sender won't overflow receiver's buffer by transmitting too much, too fast



- speed-matching service: matching the send rate to the receiving app's drain rate

- app process may be slow at reading from buffer



TCP Flow control

- Host sends 1 byte in a segment
(Inefficient usage of bandwidth)
 - 40 bytes TCP+IP headers
 - 1 byte application data

The “Silly Window Syndrome”

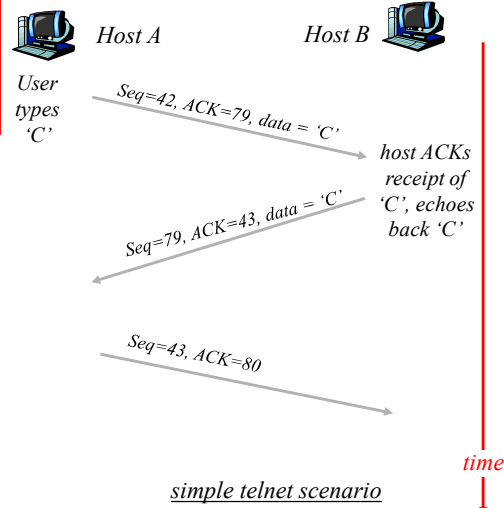
Receiver: opens window of 1 byte;

Sender: sends one-byte segment

Problem: Extremely inefficient

Solutions:

- after advertising window of size 0, receiver can only open MSS-size window
- Nagle’s Algorithm



simple telnet scenario

Nagle’s Algorithm

Nagle’s Algorithm

```

if available data and window  $\geq$  MSS:
    send a full segment
else
    if there is unACKed data in transmit
        buffer the new data until ACK received
    else
        send all the new data
    
```

Note: May result in a single byte per RTT; to turn off Nagle’s algorithm, use TCP_NODELAY option

TCP Connection Management

- Connection Establishment

Initialize TCP variables:

- seq. #s, buffers, flow control info (e.g. **RcvWindow**)

Three way handshake:

Step 1: client host sends TCP SYN segment to server


- specifies initial seq #
- no data

Step 2: server host receives SYN, replies with SYN+ACK segment

- server allocates buffers
- specifies server initial seq. #

Step 3: client receives SYN+ACK, replies with ACK segment, which may contain data

Q: How to choose the initial sequence number?



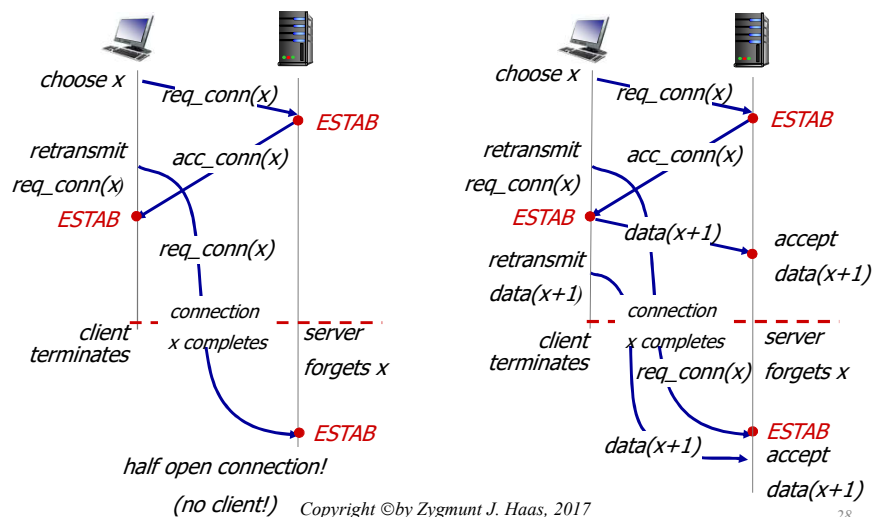
Wireless Networks Laboratory

Copyright ©by Zygmunt J. Haas, 2017

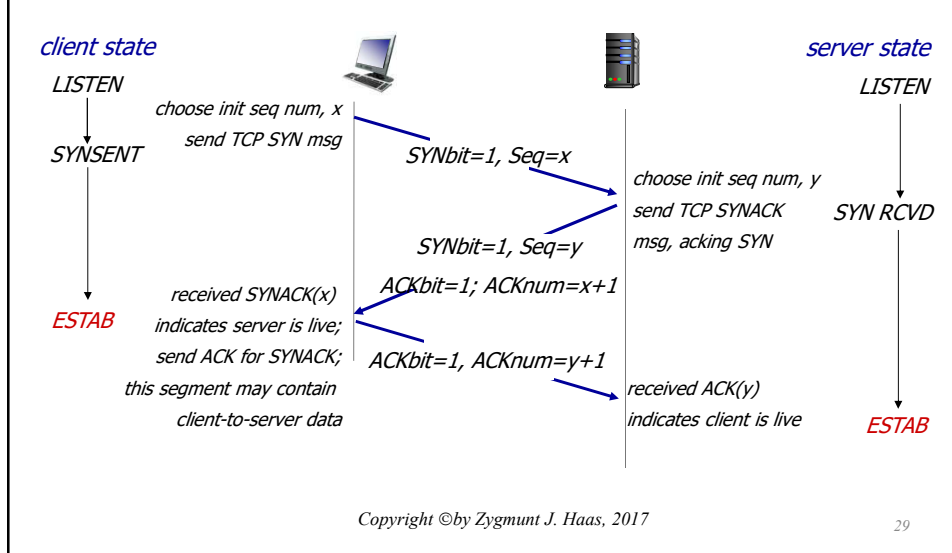
27

Agreeing to Establish a Connection

2-way handshake failure scenarios:



TCP 3-way Handshake



29

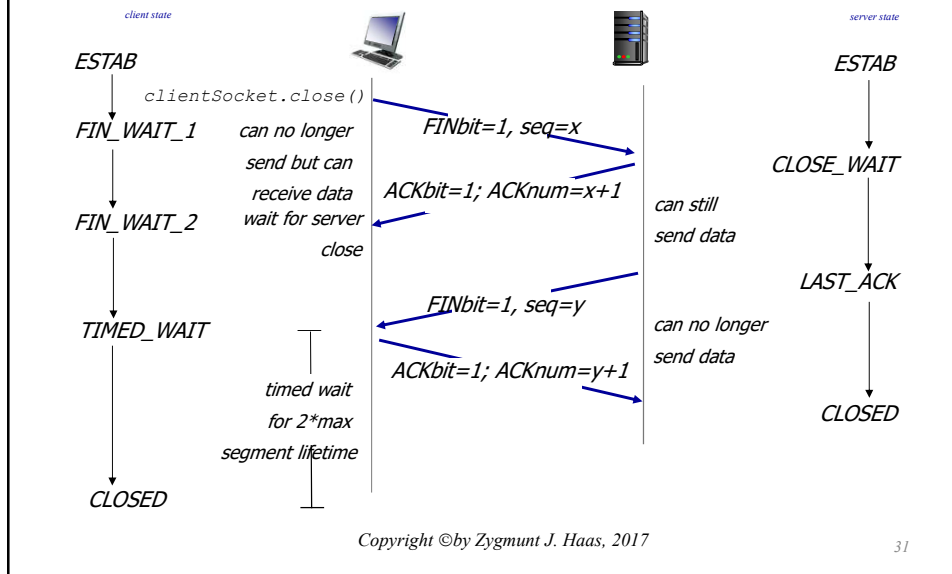
TCP: Closing a Connection

- ❖ client, server each close their side of connection
 - send TCP segment with FIN bit = 1
- ❖ respond to received FIN with ACK
 - on receiving FIN, ACK can be combined with own FIN
- ❖ simultaneous FIN exchanges can be handled

Copyright ©by Zygmunt J. Haas, 2017

30

TCP: Closing a Connection




TCP Connection Management (cont.)

Step 3: client receives *FIN*, replies with *ACK*.

- Enters “timed wait” - will respond with *ACK* to received *FIN*s

Step 4: server, receives *ACK*.
Connection closed.



Protection Against Wrap Around

- 32-bit **SequenceNum**

Bandwidth	Time Until Wrap Around
T1 (1.5 Mbps)	6.4 hours
Ethernet (10 Mbps)	57 minutes
T3 (45 Mbps)	13 minutes
FDDI (100 Mbps)	6 minutes
STS-3 (155 Mbps)	4 minutes
STS-12 (622 Mbps)	55 seconds
STS-24 (1.2 Gbps)	28 seconds

- Implication – Reduced throughput
 - MSL (Maximum Segment Lifetime) assumed to be 120 secs
 - No two TCP segments can have the same SeqNo within 120 secs

Performance: Keeping the Pipe Full

- 16-bit **AdvertisedWindow**

Bandwidth	Delay x Bandwidth Product
T1 (1.5 Mbps)	18KB (= $1.5 \times 10^6 / 8 \times .1$)
Ethernet (10 Mbps)	125KB
T3 (45 Mbps)	562.5KB
FDDI (100 Mbps)	1.2MB
STS-3 (155 Mbps)	1.9MB
STS-12 (622 Mbps)	7.7MB
STS-24 (1.2 Gbps)	15MB

The amount of in-transit unACKed TCP data

- Delay = “round-trip time” (100 msec in the above example)
- 16 bit receiver window represents at most: $2^{16} \times 8 = 0.5$ MB of data